Prompt 1: Are these reviews positive or negative?

Response 1:

The reviews are mixed, with a slight lean negative. I read all 110 and sorted them by overall sentiment:

Sentiment	Count	Share
Negative	39	35%
Positive	35	32%
Mixed (real praise and real complaints)	21	19%
Neutral, questions, or shipping-only comments	15	14%

- **Sarcasm:** 10 of the 39 negatives are sarcastic and would read as positive to a basic keyword-based sentiment tool. Examples are “Love that the ‘waterproof’ pack turned my sleeping bag into a sponge. Ten stars” and “Wow, they lasted a whole three days before disconnecting every five minutes.”
- **Products:** Each of the five products has 22 reviews, and none stands out as clearly good or bad. The Lumen Smart Bulbs are the weakest (9 negative vs. 6 positive), mostly because of app, Wi-Fi, and reliability problems. The Cozy Knit Throw Blanket is the strongest (8 positive vs. 6 negative).
- **Complaints:** Durability failures (broken zippers, buckles, switches, lids, and stopped charging) and listing claims that don’t match the product (waterproof, noise cancelling, size) come up often.
- **Praise:** Customers mostly like value for the price, comfort, design, and good customer service when it responds.
- **Customer service:** Experiences vary. Two Hearth Electric Kettle reviews and one Lumen Smart Bulbs review describe quick fixes, while one Breeze Wireless Earbuds review reports no reply at all.

Some of the 21 mixed reviews could go either way, so the exact counts shift a few points depending on how you treat them. Would you like me to add a sentiment column to the CSV, or break the complaints down by product?

Prompt 2: You are a data analyst for a global shipping company, and must sort all reviews into either positive (customer overall happy), negative (customer overall unhappy) or mixed (no clear bias towards positive/negative feelings). Return all reviews in a sorted list, with each review being labeled into its given outcome and product.

Response 2:

All 110 reviews are sorted and labeled in **sorted_reviews.csv**. They are grouped as Positive, then Mixed, then Negative, and by product within each group. The file has review_id, product, sentiment and review_text columns.

Totals: 40 Positive, 30 Mixed, 40 Negative

Product	Positive	Mixed	Negative
Trailmark Hiking Backpack	9	5	8
Cozy Knit Throw Blanket	9	7	6
Hearth Electric Kettle	8	5	9
Breeze Wireless Earbuds	7	7	8
Lumen Smart Bulbs (4-Pack)	7	6	9

How I handled the judgment calls:

- **Sarcasm** is labeled by the real sentiment, so it counts as Negative. Examples are R022 (“waterproof” pack turned sleeping bag into a sponge), R030, R066, R068, R078, R080, R085, R098 and R108.
- **Questions with no opinion** (R012, R067, R087, R101) and **neutral statements** (R026, R037, R089, R094, R104) are labeled Mixed, since there is no clear bias either way.
- **Shipping or delivery-only complaints** (R057, R062, R093) are Negative, because the customer was clearly unhappy. When the product itself turned out fine, I leaned Positive, as in R014 and R056.
- **A problem that was resolved** is Positive when the customer ended up happy (R084). R013 is Mixed, because the replacement “works so far” but the first unit failed.
- **Praise outweighed by a minor gripe** is Positive (R059, R105). Roughly balanced praise and complaint is Mixed (R055, R060, R071, R097).

If you’d rather have the neutral and question-only reviews in a separate “Neutral/No sentiment” category instead of Mixed, I can split them out.

Snippet of all labeled reviews (this is the only response where it showed every review labeled):

R006	Breeze Wireless Earbuds	Positive	Perfect for my commute. Clear calls, solid battery.

R033	Breeze Wireless Earbuds	Positive	Best purchase I've made all year. My old $200 pair is collecting dust now.
R052	Breeze Wireless Earbuds	Positive	I was skeptical but these are fantastic. Would buy again.
R054	Breeze Wireless Earbuds	Positive	Super comfortable, I forget I'm wearing them. Pairing was instant.
R056	Breeze Wireless Earbuds	Positive	Box arrived crushed but the earbuds inside were fine. They work well.
R072	Breeze Wireless Earbuds	Positive	Not as good as my AirPods but for a quarter of the price I can't complain.
R077	Breeze Wireless Earbuds	Positive	Sound quality is amazing for the price. Bass is punchy and they stay in during runs.
R036	Cozy Knit Throw Blanket	Positive	Ordered the gray and got navy. Kept it anyway since it's nice.
R040	Cozy Knit Throw Blanket	Positive	My toddler drags it everywhere and it still looks new. Impressive.
R045	Cozy Knit Throw Blanket	Positive	Perfect weight for summer evenings on the porch.
R046	Cozy Knit Throw Blanket	Positive	So soft and warm. My cat has claimed it as hers.
R069	Cozy Knit Throw Blanket	Positive	love love love!!! buying another
R079	Cozy Knit Throw Blanket	Positive	Gave one to each of my daughters for Christmas. They both love them.
R081	Cozy Knit Throw Blanket	Positive	Washes well, no pilling after many cycles. Great quality.
R083	Cozy Knit Throw Blanket	Positive	Looks way more expensive than it is. Guests always compliment it.
R102	Cozy Knit Throw Blanket	Positive	Best throw blanket I've owned. Heavy, warm, and soft.

Prompt 3: Are these reviews positive or negative overall? For example, this review: “R009,Lumen Smart Bulbs (4-Pack),"Great bulbs, terrible app. If you can get past setup you'll

love them." This is a positive review despite some negative sentiment not targeted at the product itself. And this review: “R026,Lumen Smart Bulbs (4-Pack),Package arrived a day early.” Can be seen as positive, as it is showcasing a positive outcome despite a lack of emotion.

Response 3:

Overall, the reviews are split almost evenly, with a very slight lean positive. Of the 110 reviews, I classified them like this:

Sentiment	Count	Share
Positive	43	39%
Negative	40	36%
Mixed	15	14%
Neutral	8	7%
Question (no sentiment)	4	4%

Among reviews that take a side, that’s about 52% positive to 48% negative. Most of the 15 mixed reviews lean positive (like R008 and R055), and a few lean negative (like R073 and R099). So the real picture is a bit warmer than the raw counts suggest.

I followed your rubric. I counted a review as positive when the product itself came out fine, even if something peripheral went wrong (R009). I also counted it as positive when it described a good outcome without any emotion (R026). That is why I rated these as positive:

- **R014:** the kettle arrived dented, but customer service fixed it and the kettle is “fine”.
- **R056:** the box was crushed, but the earbuds “work well”.
- **R084:** pairing was a problem, but support solved it.
- **R036:** the wrong color arrived, but “kept it anyway since it’s nice”.

Sarcasm was the biggest trap. About 11 reviews read as positive on the surface but are negative, such as R022 (“Ten stars” for a “waterproof” pack that soaked the sleeping bag), R066 (“Great feature” for earbuds that fall out) and R108 (“Truly impressive engineering”). I scored these as negative. A keyword-based sentiment tool would likely get them wrong.

Judgment calls:

- **Delivery problems:** I rated reviews about delivery or fulfillment problems as negative when the outcome was bad (R062 “Shipping took 3 weeks”, R093 “left it in the rain”). That mirrors your R026 logic in reverse.

- **Product vs. app:** I treated complaints about the product as negative even when the cause was software, such as the app or the account/location requirement (R096). R009 differs because the reviewer still recommends the bulbs.
- **Neutral reviews:** I kept flat, functional reviews like R037 (“They light up”) and R043 (“It’s fine”) as neutral. Under your R026 logic you could lean them positive, which would push the positive share to about 46%.

By product, each had 22 reviews:

Product	Positive	Negative	Mixed	Neutral	Question
Trailmark Hiking Backpack	10	8	2	1	1
Cozy Knit Throw Blanket	9	6	4	2	1
Lumen Smart Bulbs (4-Pack)	9	9	2	1	1
Hearth Electric Kettle	8	9	3	2	0
Breeze Wireless Earbuds	7	8	4	2	1

The throw blanket and backpack look best. The earbuds are weakest, with the most mixed reviews and several complaints about durability and fit.

I read and labeled each review by hand rather than using a scoring tool. If you want, I can export the full review-by-review labels as a CSV so you can check my calls.

Prompt 4: Are these reviews positive or negative? Only showcase how many reviews are classified as positive and negative, removing mixed reviews. Please only include the final answer in your response.

Response 4:

Positive: 34 reviews
Negative: 39 reviews

Mixed, neutral and question-only reviews were excluded. Sarcastic reviews were counted by their actual sentiment, so “Love that the waterproof pack turned my sleeping bag into a sponge” counts as negative.

Spot-check Table:

R009,Lumen Smart Bulbs (4-Pack),Great bulbs, terrible app. If you can get past setup you'll love them. - Positive
R026,Lumen Smart Bulbs (4-Pack),Package arrived a day early. - Positive
R001,Trailmark Hiking Backpack,Smells like chemicals out of the box. Took a week to air out. -

Negative
R042,Cozy Knit Throw Blanket,Ok for the price - Positive
R043,Hearth Electric Kettle,"It's fine. Does the job, nothing special." - Positive
R068,Cozy Knit Throw Blanket,"Wonderful, it's itchy AND thin. Two for one." - Negative
R072,Breeze Wireless Earbuds,Not as good as my AirPods but for a quarter of the price I can't complain. - Positive
R073,Lumen Smart Bulbs (4-Pack),Good enough for the price. I'd buy a better brand next time though. - Positive
R104,Hearth Electric Kettle,It boils water. It's a kettle. - Positive
R108,Breeze Wireless Earbuds,"Wow, they lasted a whole three days before disconnecting every five minutes. Truly impressive engineering." - Negative

Compared to the best prompt, we agreed on R072, R001, R068, and R108, but disagreed on R026, R042, R043, R009, R073, and R104. I believe my decisions are more correct, as almost every disagreement stems from the LLM categorizing reviews that were clearly positive or negative as mixed, due to either unenthusiastic reviews or negative sentiment towards minor aspects.

Analysis:

I used role prompting in my second prompt, where I said the LLM was a data analyst. I used constraints in prompt 4, with me constraining the LLM to only showing the final tally without needing to ‘show its work’. I used few-shot prompting in prompt 3, giving an example of some tricky reviews that were positive. I only specified output format in prompt 4 and 2, however Claude ended up having the same output format across every prompt except 4, where I specifically was trying to have a response in a different format. I feel the technique that ended up resulting in better prompts was only few-shot prompting, as every other change between my prompts resulted in almost no real changes in the output of the LLM. I think this is because of two reasons, first Claude being a more coding centric model is more structured in its responses overall, and due to that would end up implicitly doing output formatting and categorization on its own. Secondly few-shot prompting allowed for Claude to get help on more mixed reviews in order to better point them in one way or another rather than simply making them mixed. I noticed the LLM would focus very hard on putting any review with multiple tones in mixed, when I feel they could have been better categorized into positive/negative. Interestingly, only one of my prompts resulted in Claude returning a list of every review and their labels, most likely due to me stating it needed to return all of the labeled reviews only in this prompt.